

Diabetes Dataset Testing

2023-09-24

Functions

```
count_na_per_column <- function(df) {
  sapply(df, function(x) sum(is.na(x)))
}

# Create a function to calculate metrics
calc_metrics <- function(pred, true) {
  confusion <- table(pred, true)
  TP <- confusion[2, 2]
  FP <- confusion[2, 1]
  TN <- confusion[1, 1]
  FN <- confusion[1, 2]
  Sensitivity <- TP / (TP + FN)
  Specificity <- TN / (TN + FP)
  Accuracy <- (TP + TN) / (TP + FP + TN + FN)
  Precision <- TP / (TP + FP)
  F1 <- 2 * (Precision * Sensitivity) / (Precision + Sensitivity)
  pred_obj <- prediction(as.numeric(pred), as.numeric(true))
  perf <- performance(pred_obj, "auc")
  AUC <- as.numeric(perf@y.values)
  return(c(AUC, Sensitivity, Specificity, F1, Accuracy))
}

### A string like "Logistic" has to be put in. <--- from the colnames of df_coef
importance_plot <- function(model_string) {

  # Create a data frame for plotting
  coef_df <- data.frame(
    Variable = rownames(df_coef)[-1], # Exclude the intercept
    Importance = abs(df_coef[-1, model_string]) # Exclude the intercept
  )

  # Order the variables by importance
  coef_df <- coef_df[order(coef_df$Importance), ]

  # Plot
  ggplot(coef_df, aes(x = reorder(Variable, Importance), y = Importance, fill = Importance)) +
    geom_bar(stat = "identity") +
    scale_fill_gradient(low = "lightblue2", high = "lightblue3") +
    theme_minimal() +
    theme(axis.text.y = element_text(size = 12),
          axis.text.x = element_text(size = 10),
          title = element_text(size = 15),
```

```

    axis.title.y = element_blank(),
    axis.title.x = element_blank(),
    legend.position = "none" +
coord_flip() +
ggtitle(paste(model_string, "Model - Variable Importance"))
}

```

5. Data Processing

```

# Loading the data
data("PimaIndiansDiabetes2", package = "mlbench")
diabetes <- PimaIndiansDiabetes2
diabetes$diabetes <- as.factor(ifelse(diabetes$diabetes == "pos", 1, 0))
#diabetes$diabetes <- ifelse(diabetes$diabetes == "pos", 1, 0)
#kable(t(count_na_per_column(diabetes)))

```

Removing NA's

```

diabetes <- na.omit(diabetes)
#kable(t(count_na_per_column(diabetes)))

```

Removing Outliers

```

outliers <- check_outliers(diabetes, method = "mahalanobis")
#plot(outliers)
outliers <- as.vector(outliers)

#diabetes <- diabetes[!outliers, ]

```

Data for models

```

# Split
set.seed(222)
n <- nrow(diabetes)
training.samples <- sample(1:n, size = 0.75 * n)
train.data <- diabetes[training.samples, ]
scaled_train.data <- scale(train.data[, 1:8])
train.data[, 1:8] <- scaled_train.data
#train.data <- smote(diabetes ~ ., train.data, perc.over = 1)
test.data <- diabetes[-training.samples, ]
scaled_test.data <- scale(test.data[, 1:8])
test.data[, 1:8] <- scaled_test.data
#test.data <- na.omit(test.data)

# Handle NA's in training set
#mice <- complete(mice(subset(train.data, select = -c(triceps, insulin)), method='rf', seed = 123))
#mice <- complete(mice(train.data, method='rf', seed = 123))

#train.data$glucose <- mice$glucose

```

```

#train.data$pressure <- mice$pressure
#train.data$mass <- mice$mass

#train.data <- na.omit(train.data)

#mice.triceps <- complete(mice(subset(train.data, select = -insulin), method='rf', seed = 123))
#train.data$triceps <- mice.triceps$triceps

#mice.insulin <- complete(mice(train.data, method='rf', seed = 123))
#train.data$insulin <- mice.insulin$insulin

# Train
X.train <- model.matrix(diabetes~., data = train.data)[,-1]
#X.train <- scale(X.train)
y.train <- train.data$diabetes

# Test
X.test <- model.matrix(diabetes ~., data = test.data)[,-1]
#X.test <- scale(X.test)
y.test <- test.data$diabetes

```

Logistic Regression

```

set.seed(123)
log.model <- glm(diabetes ~., data = train.data, family = "binomial")

# Make predictions
probabilities <- predict(log.model, newdata = test.data, type = "response")
predicted.classes.log <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
# mean(predicted.classes == y.test)
log.conf <- confusionMatrix(predicted.classes.log, y.test, positive = "1")
log.conf

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 61 10
##           1  6 21
##
##           Accuracy : 0.8367
##           95% CI : (0.7484, 0.9037)
##           No Information Rate : 0.6837
##           P-Value [Acc > NIR] : 0.0004538
##
##           Kappa : 0.609
##
##           McNemar's Test P-Value : 0.4532547
##
##           Sensitivity : 0.6774
##           Specificity : 0.9104

```

```
##          Pos Pred Value : 0.7778
##          Neg Pred Value : 0.8592
##          Prevalence : 0.3163
##          Detection Rate : 0.2143
##          Detection Prevalence : 0.2755
##          Balanced Accuracy : 0.7939
##
##          'Positive' Class : 1
##
```

LASSO

```
set.seed(123)
cv.lasso.model <- cv.glmnet(X.train, y.train, alpha = 1, family = "binomial",
                           intercept = T)
#plot(cv.lasso.model)
cbind(coef(cv.lasso.model, s = cv.lasso.model$lambda.min), coef(cv.lasso.model, s = cv.lasso.model$lambda.1se))

## 9 x 2 sparse Matrix of class "dgCMatrix"
##          s1          s1
## (Intercept) -0.9087548 -0.79367862
## pregnant    0.2500270  0.04713706
## glucose     1.0571306  0.80577102
## pressure     .          .
## triceps     0.0626410  .
## insulin      .          .
## mass        0.4184043  0.19426613
## pedigree    0.3622318  0.10361089
## age         0.2235214  0.17759602

# Make predictions
probabilities <- predict(cv.lasso.model, newx = X.test, s = cv.lasso.model$lambda.min, type = "response")
predicted.classes.lasso.min <- as.factor(ifelse(probabilities > 0.5, 1, 0))

probabilities <- predict(cv.lasso.model, newx = X.test, s = cv.lasso.model$lambda.1se, type = "response")
predicted.classes.lasso.1se <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
#mean(predicted.classes.lasso.min == y.test)
lasso.min.conf <- confusionMatrix(predicted.classes.lasso.min, y.test, positive = "1")
lasso.min.conf

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 61 10
##          1  6 21
##
##          Accuracy : 0.8367
##          95% CI : (0.7484, 0.9037)
##          No Information Rate : 0.6837
##          P-Value [Acc > NIR] : 0.0004538
##
##          Kappa : 0.609
```

```
##
## McNemar's Test P-Value : 0.4532547
##
##          Sensitivity : 0.6774
##          Specificity : 0.9104
##          Pos Pred Value : 0.7778
##          Neg Pred Value : 0.8592
##          Prevalence : 0.3163
##          Detection Rate : 0.2143
##          Detection Prevalence : 0.2755
##          Balanced Accuracy : 0.7939
##
##          'Positive' Class : 1
##
lasso.1se.conf <- confusionMatrix(predicted.classes.lasso.1se, y.test, positive = "1")
lasso.1se.conf
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 62 12
##          1  5 19
##
##          Accuracy : 0.8265
##          95% CI : (0.7369, 0.8956)
##          No Information Rate : 0.6837
##          P-Value [Acc > NIR] : 0.001064
##
##          Kappa : 0.573
##
## McNemar's Test P-Value : 0.145610
##
##          Sensitivity : 0.6129
##          Specificity : 0.9254
##          Pos Pred Value : 0.7917
##          Neg Pred Value : 0.8378
##          Prevalence : 0.3163
##          Detection Rate : 0.1939
##          Detection Prevalence : 0.2449
##          Balanced Accuracy : 0.7691
##
##          'Positive' Class : 1
##
```

in data: set.seed(123) and train 0.75 all others see(123) i like it a lot. min and 1se very sparse and good pred, just not much difference in amount of sparsity and amount of FN

in data: set.seed(2) and train 0.75 all others see(123) i like it a lot. 1se sparse and more FN but same pred

in data: set.seed(222) and train 0.75 all others see(123) pretty good similar too above

in data: set.seed(6) and train 0.75 all others see(123) gives very sparse 1se lasso and okay pred

in data: set.seed(13) and train 0.75 all others see(123) pretty nice! sparse 1se, but pred is same

in data: set.seed(42) and train 0.75 all others see(123) pretty nice! very very sparse 1se, more FN in 1se

some things we see:

- the more conservative model (1se) usually has more FN —> it makes the safe/conservative choice of going for the class that has 2/3 of observations

Ridge

```
set.seed(123)
cv.ridge.model <- cv.glmnet(X.train, y.train, alpha = 0, family = "binomial", intercept = T)
#plot(cv.ridge)

# Make predictions
probabilities <- predict(cv.ridge.model, newx = X.test, s = cv.ridge.model$lambda.min, type = "response")
predicted.classes.ridge <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
#mean(predicted.classes.ridge == y.test)
ridge.conf <- confusionMatrix(as.factor(predicted.classes.ridge), y.test, positive = "1")
ridge.conf

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##          0 62 10
##          1  5 21
##
##              Accuracy : 0.8469
##              95% CI : (0.7601, 0.9117)
##      No Information Rate : 0.6837
##      P-Value [Acc > NIR] : 0.0001806
##
##              Kappa : 0.6301
##
##  Mcnemar's Test P-Value : 0.3016996
##
##              Sensitivity : 0.6774
##              Specificity : 0.9254
##      Pos Pred Value : 0.8077
##      Neg Pred Value : 0.8611
##              Prevalence : 0.3163
##      Detection Rate : 0.2143
##      Detection Prevalence : 0.2653
##      Balanced Accuracy : 0.8014
##
##      'Positive' Class : 1
##
```

Elastic Net

```
set.seed(123)

# CV and tuning grid
```

```

myFolds <- createFolds(y.train, k = 10, list = TRUE)
myControl <- trainControl(index = myFolds)

#cvControl <- trainControl(method = "cv", number = 10)
tuneGrid <- expand.grid(alpha = seq(0, 1, by = 0.05), lambda = 10^seq(1, -3, length=100))

# Train model
elasticnet.model <- train(X.train, y.train, method = "glmnet", trControl = myControl,
                          tuneGrid = tuneGrid, intercept = T, family = "binomial")

# Make predictions
probabilities <- predict(elasticnet.model, newdata = X.test, type = "prob") # "raw" already outputs pre
predicted.classes.elasticnet <- ifelse(probabilities$`1` > 0.5, 1, 0)

# Accuracy
# mean(predicted.classes == y.test)
elasticnet.conf <- confusionMatrix(as.factor(predicted.classes.elasticnet), y.test, positive = "1")
elasticnet.conf

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 63 11
##           1  4 20
##
##           Accuracy : 0.8469
##           95% CI : (0.7601, 0.9117)
##    No Information Rate : 0.6837
##    P-Value [Acc > NIR] : 0.0001806
##
##           Kappa : 0.6233
##
##  McNemar's Test P-Value : 0.1213353
##
##           Sensitivity : 0.6452
##           Specificity : 0.9403
##           Pos Pred Value : 0.8333
##           Neg Pred Value : 0.8514
##           Prevalence : 0.3163
##           Detection Rate : 0.2041
##           Detection Prevalence : 0.2449
##           Balanced Accuracy : 0.7927
##
##           'Positive' Class : 1
##
# Output coefficients
coef(elasticnet.model$finalModel, s = elasticnet.model$bestTune$lambda)


## 9 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -0.783644314
## pregnant    0.124519132
## glucose      0.616288277

```

```
## pressure      0.003625147
## triceps       0.073458775
## insulin       0.067361284
## mass          0.200698201
## pedigree      0.158043114
## age           0.188599082

# Filter data to only include specific alpha values
filtered_data <- subset(elasticnet.model$results, alpha %in% seq(0, 1, by=0.1))

# Create the plot
ggplot(filtered_data, aes(x=lambda, y=Accuracy, color=factor(alpha))) +
  geom_point() +
  geom_line() +
  scale_x_log10() +
  scale_color_discrete(name = "Alpha") +
  labs(x = "Lambda", y = "Accuracy") +
  theme_minimal()
```




```
library(ggcorrplot)
# Ensure that all remaining columns after subsetting are numeric
numeric_data <- subset(diabetes, select = -c(diabetes))
corr <- round(cor(numeric_data), 1)
ggcorrplot(corr,
  type = "lower",
  lab = TRUE,
  lab_size = 3,
```



```

colors = c("red", "white", "cyan4"),
title = "Correlogram of Diabetes Dataset",
ggtheme = theme_bw())

```



```


result <- check_collinearity(log.model)
plot(result)

```

```
## Variable 'Component' is not in your data frame :/
```


Collinearity

High collinearity (VIF) may inflate parameter uncertainty



Outliers

```
result <- check_outliers(diabetes, method = "mahalanobis")  
plot(result, type = "dots")
```



```
result
```

```
## 13 outliers detected: cases 2, 4, 5, 58, 78, 87, 111, 120, 185, 226,
## 236, 251, 302.
## - Based on the following method and threshold: mahalanobis (30).
## - For variables: pregnant, glucose, pressure, triceps, insulin, mass,
## pedigree, age.
```

```
as.vector(result)
```

```
## [1] FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [37] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [49] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
## [61] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [73] FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE
## [85] FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [97] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [109] FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [121] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [133] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [145] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [169] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [181] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
## [193] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
## [205] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [217] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE
## [229] FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
## [241] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
## [253] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [265] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [277] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [289] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [301] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [313] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [325] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [337] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [349] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [361] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [373] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [385] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

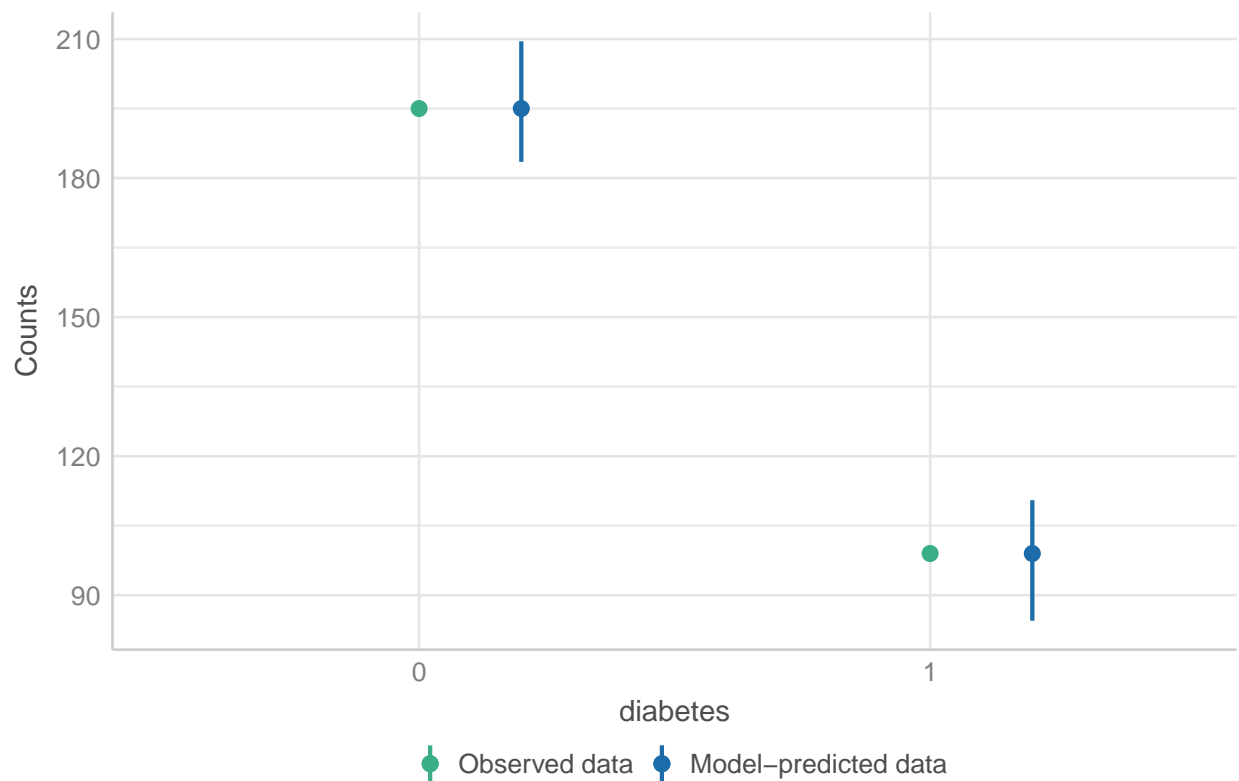
Check model

```
check <- check_model(log.model, panel = F)
plot(check)
```

```
## $PP_CHECK
```


Posterior Predictive Check

Model-predicted intervals should include observed data points



```
##
## $BINNED_RESID
```


Binned Residuals
Points should be within error bounds



\$OUTLIERS

Influential Observations


Points should be inside the contour lines



\$VIF

Collinearity


High collinearity (VIF) may inflate parameter uncertainty



##

\$QQ

Normality of Residuals
Dots should fall along the line



```
binmed_residuals(log.model)
```

```
## Warning: About 88% of the residuals are inside the error bounds (~95% or higher would be good).
```

Calculate model metrics

```
# Calculate metrics for each model
metrics.log <- calc_metrics(predicted.classes.log, y.test)
metrics.lasso.min <- calc_metrics(predicted.classes.lasso.min, y.test)
metrics.lasso.1se <- calc_metrics(predicted.classes.lasso.1se, y.test)
metrics.ridge <- calc_metrics(predicted.classes.ridge, y.test)
metrics.elasticnet <- calc_metrics(predicted.classes.elasticnet, y.test)

# Store in a dataframe
df_metrics <- as.data.frame(matrix(c(metrics.log, metrics.lasso.min, metrics.lasso.1se, metrics.ridge, metrics.elasticnet),
  nrow = 5, ncol = 5))
colnames(df_metrics) <- c("Logistic", "Lasso_min", "Lasso_1se", "Ridge", "ElasticNet")
rownames(df_metrics) <- c("AUC", "Sensitivity", "Specificity", "F1", "Accuracy")
df_metrics
```

```
##           Logistic Lasso_min Lasso_1se      Ridge ElasticNet
## AUC          0.7939336 0.7939336 0.7691382 0.8013962 0.7927299
## Sensitivity 0.6774194 0.6774194 0.6129032 0.6774194 0.6451613
## Specificity 0.9104478 0.9104478 0.9253731 0.9253731 0.9402985
## F1           0.7241379 0.7241379 0.6909091 0.7368421 0.7272727
## Accuracy    0.8367347 0.8367347 0.8265306 0.8469388 0.8469388
```


Latex Table Metrics

```
metrics_latex <- df_metrics %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), digits = 3, escape = FALSE,
        col.names = c("Logistic", "Lasso  $\\lambda_{\\text{min}}$ ",
                      "Lasso  $\\lambda_{\\text{1se}}$ ", "Ridge", "ElasticNet")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Models" = 5), bold = TRUE) %>%
  row_spec(nrow(df_metrics) - 1, extra_latex_after = "\\midrule[.08em]") %>%
  row_spec(nrow(df_metrics), bold = TRUE)

print(metrics_latex)
```

	Models				
	Logistic	Lasso λ_{\min}	Lasso λ_{1se}	Ridge	ElasticNet
AUC	0.794	0.794	0.769	0.801	0.793
Sensitivity	0.677	0.677	0.613	0.677	0.645
Specificity	0.910	0.910	0.925	0.925	0.940
F1	0.724	0.724	0.691	0.737	0.727
Accuracy	0.837	0.837	0.827	0.847	0.847

Data Frame of all model coefficients

```
coef.log <- coef(log.model)
coef.lasso.min <- (coef(cv.lasso.model, s = cv.lasso.model$lambda.min))
coef.lasso.1se <- coef(cv.lasso.model, s = cv.lasso.model$lambda.1se)
coef.ridge <- coef(cv.ridge.model)
coef.elasticnet <- coef(elasticnet.model$finalModel, s = elasticnet.model$bestTune$lambda)

# Create a data frame to store coefficients
df_coef <- data.frame(
  Logistic = as.vector(coef.log),
  Lasso_min = as.vector(as.matrix(coef.lasso.min)),
  Lasso_1se = as.vector(as.matrix(coef.lasso.1se)),
  Ridge = as.vector(as.matrix(coef.ridge)),
  ElasticNet = as.vector(as.matrix(coef.elasticnet))
)
rownames(df_coef) <- names(coef.log)

options(scipen = 999)
df_coef_sparse <- round(df_coef, 4)
df_coef_sparse[] <- apply(df_coef, 2, function(x) ifelse(x == 0, '.', x))

df_coef_sparse

##                Logistic                Lasso_min                Lasso_1se
## (Intercept) -0.976333973381479 -0.908754819703348 -0.793678618299817
```

```
## pregnant      0.348383290980628    0.25002699073021 0.0471370598584918
## glucose       1.19662756034941    1.05713059005187 0.805771016005861
## pressure     -0.00787072837112941      .
## triceps       0.0990151553153173 0.0626409984004955      .
## insulin      -0.0108165468422877      .
## mass          0.527217247967629    0.418404324962646    0.19426613242444
## pedigree     0.488903494680658    0.362231763229765    0.103610887531999
## age          0.24003005800779    0.22352140524853    0.177596021359124
##              Ridge      ElasticNet
## (Intercept) -0.776920592058855    -0.783644314126479
## pregnant    0.156750171833101    0.124519132142544
## glucose     0.437363939996492    0.616288276600248
## pressure    0.084266432087758    0.00362514721615527
## triceps     0.121285665042735    0.0734587750905991
## insulin     0.157046443194511    0.0673612838310804
## mass        0.181212025582025    0.200698200930468
## pedigree    0.176979262379118    0.158043114063933
## age         0.186020503416058    0.1885990817328
```

Latex Table Coef

```
coef_latex <- df_coef_sparse %>%
  kable(digits = 3, format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Logistic", "Lasso  $\lambda_{\min}$ ",
                      "Lasso  $\lambda_{1se}$ ", "Ridge", "ElasticNet")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Models" = 5), bold = TRUE)

print(coef_latex)
```

	Models				
	Logistic	Lasso λ_{\min}	Lasso λ_{1se}	Ridge	ElasticNet
(Intercept)	-0.976333973381479	-0.908754819703348	-0.793678618299817	-0.776920592058855	-0.783644314126479
pregnant	0.348383290980628	0.25002699073021	0.0471370598584918	0.156750171833101	0.124519132142544
glucose	1.19662756034941	1.05713059005187	0.805771016005861	0.437363939996492	0.616288276600248
pressure	-0.00787072837112941	.	.	0.084266432087758	0.00362514721615527
triceps	0.0990151553153173	0.0626409984004955	.	0.121285665042735	0.0734587750905991
insulin	-0.0108165468422877	.	.	0.157046443194511	0.0673612838310804
mass	0.527217247967629	0.418404324962646	0.19426613242444	0.181212025582025	0.200698200930468
pedigree	0.488903494680658	0.362231763229765	0.103610887531999	0.176979262379118	0.158043114063933
age	0.24003005800779	0.22352140524853	0.177596021359124	0.186020503416058	0.1885990817328

Confusion matrix latex - log.model

```
# Convert the table to a data frame
log.conf.df <- as.data.frame(log.conf$table)


# Replace the numeric labels with text labels
log.conf.df$Prediction <- factor(log.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
log.conf.df$Reference <- factor(log.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))

# Rename columns as required by cums
```

```
names(log.conf.df) <- c('Target', 'Prediction', 'N')
# Plot the confusion matrix using cums
plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
  add_row_percentages = FALSE, add_col_percentages = FALSE, rotate_y_text = TRUE,
  place_x_axis_above = TRUE, class_order = c("Positive", "Negative"))
```

```
## Warning in plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
## add_row_percentages = FALSE, : 'ggimage' is missing. Will not plot arrows and
## zero-shading.
```

```
## Warning in plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
## add_row_percentages = FALSE, : 'rsvg' is missing. Will not plot arrows and
## zero-shading.
```




```
### Confusion matrix latex - lasso.min.model
```

```
# For lasso.min.conf
lasso.min.conf.df <- as.data.frame(lasso.min.conf$table)
lasso.min.conf.df$Prediction <- factor(lasso.min.conf.df$Prediction, levels = c("0", "1"), labels = c("N", "P"))
lasso.min.conf.df$Reference <- factor(lasso.min.conf.df$Reference, levels = c("0", "1"), labels = c("N", "P"))
names(lasso.min.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_p
```

```
## Warning in plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.
```

```
## Warning in plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - lasso.1se.model

```
lasso.1se.conf.df <- as.data.frame(lasso.1se.conf$table)
lasso.1se.conf.df$Prediction <- factor(lasso.1se.conf.df$Prediction, levels = c("0", "1"), labels = c("Neg", "Pos"))
lasso.1se.conf.df$Reference <- factor(lasso.1se.conf.df$Reference, levels = c("0", "1"), labels = c("Neg", "Pos"))
names(lasso.1se.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)

## Warning in plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.

## Warning in plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - ridge.model

```
# For ridge.conf
ridge.conf.df <- as.data.frame(ridge.conf$table)
ridge.conf.df$Prediction <- factor(ridge.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
ridge.conf.df$Reference <- factor(ridge.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))
names(ridge.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)

## Warning in plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.

## Warning in plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - elasticnet.model

```
# For elasticnet.conf
elasticnet.conf.df <- as.data.frame(elasticnet.conf$table)
elasticnet.conf.df$Prediction <- factor(elasticnet.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
elasticnet.conf.df$Reference <- factor(elasticnet.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))
names(elasticnet.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)
```

```
## Warning in plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.


## Warning in plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```



Variable Importance


```
importance_plot("Logistic")
```

Logistic Model – Variable Importance




```
importance_plot("Lasso_min")
```


Lasso_min Model – Variable Importance




```
importance_plot("Lasso_1se")
```

Lasso_1se Model – Variable Importance




```
importance_plot("Ridge")
```

Ridge Model – Variable Importance



```
importance_plot("ElasticNet")
```

ElasticNet Model – Variable Importance



Perc missing

```
library(ggplot2)
```

```
# Count of missing values per column
```

```
count_na <- c(pregnant = 0, glucose = 5, pressure = 35, triceps = 227, insulin = 374, mass = 11, pedigree = 1)
```

```
# Total number of rows in the data
```

```
n_rows <- nrow(diabetes)
```

```
# Convert to data frame and calculate percentages
```

```
df_na <- data.frame(Column = names(count_na), Count = as.numeric(count_na))
```

```
df_na$Percentage <- (df_na$Count / n_rows) * 100
```

```
# Plot
```

```
ggplot(df_na, aes(x = reorder(Column, -Percentage), y = Percentage, fill = Percentage)) +
```

```
  geom_bar(stat = "identity") +
```

```
  scale_fill_gradient(low = "lightblue1", high = "lightblue3") +
```

```
  theme_minimal() +
```

```
  theme(
```

```
    axis.text.y = element_text(size = 12),
```

```
    axis.text.x = element_text(size = 12),
```

```
    title = element_text(size = 15),
```


```
    axis.title.y = element_blank(),
```

```
    axis.title.x = element_blank(),
```

```
    legend.position = "none"
```

```
) +
```

```
ggtitle("Percentage of Missing Values per Column")
```



log model summary table latex

```
# Extract coefficients from the model summary
coef_summary <- summary(log.model)$coefficients

# Convert the matrix to a data frame for kable()
coef_summary <- round(as.data.frame(coef_summary)[, c(1, 4)], 3)
colnames(coef_summary) <- c("Estimate", "p-value")

coef_summary$`p-value` <- cell_spec(coef_summary$`p-value`, "latex",
                                   bold = ifelse(coef_summary$`p-value` < 0.05, TRUE, FALSE))

# Create the LaTeX table
kable(coef_summary, "latex", booktabs = TRUE, align = "c", escape = F) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position")) %>%
  row_spec(0, bold = TRUE)
```

	Estimate	p-value
(Intercept)	-0.976	0
pregnant	0.348	0.087
glucose	1.197	0
pressure	-0.008	0.963
triceps	0.099	0.622
insulin	-0.011	0.95
mass	0.527	0.016
pedigree	0.489	0.004
age	0.240	0.251