

Diabetes Dataset Testing

2023-09-24

Functions

```
count_na_per_column <- function(df) {
  sapply(df, function(x) sum(is.na(x)))
}

# Create a function to calculate metrics
calc_metrics <- function(pred, true) {
  confusion <- table(pred, true)
  TP <- confusion[2, 2]
  FP <- confusion[2, 1]
  TN <- confusion[1, 1]
  FN <- confusion[1, 2]
  Sensitivity <- TP / (TP + FN)
  Specificity <- TN / (TN + FP)
  Accuracy <- (TP + TN) / (TP + FP + TN + FN)
  Precision <- TP / (TP + FP)
  F1 <- 2 * (Precision * Sensitivity) / (Precision + Sensitivity)
  pred_obj <- prediction(as.numeric(pred), as.numeric(true))
  perf <- performance(pred_obj, "auc")
  AUC <- as.numeric(perf@y.values)
  return(c(AUC, Sensitivity, Precision, F1, Accuracy))
}

### A string like "Logistic" has to be put in. <--- from the colnames of df_coef
importance_plot <- function(model_string) {

  # Create a data frame for plotting
  coef_df <- data.frame(
    Variable = rownames(df_coef)[-1], # Exclude the intercept
    Importance = abs(df_coef[-1, model_string]) # Exclude the intercept
  )

  # Order the variables by importance
  coef_df <- coef_df[order(coef_df$Importance), ]

  # Plot
  ggplot(coef_df, aes(x = reorder(Variable, Importance), y = Importance, fill = Importance)) +
    geom_bar(stat = "identity") +
    scale_fill_gradient(low = "lightblue2", high = "lightblue3") +
    theme_minimal() +
    theme(axis.text.y = element_text(size = 12),
          axis.text.x = element_text(size = 10),
          title = element_text(size = 15),
```

```

    axis.title.y = element_blank(),
    axis.title.x = element_blank(),
    legend.position = "none") +
  coord_flip() +
  ggtitle("Variable Importance")
}

```

5. Data Processing

```

# Loading the data
data("PimaIndiansDiabetes2", package = "mlbench")
diabetes <- PimaIndiansDiabetes2
diabetes$diabetes <- as.factor(ifelse(diabetes$diabetes == "pos", 1, 0))
#diabetes$diabetes <- ifelse(diabetes$diabetes == "pos", 1, 0)
#kable(t(count_na_per_column(diabetes)))

```

Removing NA's

```

diabetes <- na.omit(diabetes)
#kable(t(count_na_per_column(diabetes)))

```

Removing Outliers

```

outliers <- check_outliers(diabetes, method = "mahalanobis")
#plot(outliers)
outliers <- as.vector(outliers)

#diabetes <- diabetes[!outliers, ]

```

Data for models

```

# Split
set.seed(222)
n <- nrow(diabetes)
training.samples <- sample(1:n, size = 0.75 * n)
train.data <- diabetes[training.samples, ]
scaled_train.data <- scale(train.data[, 1:8])
train.data[, 1:8] <- scaled_train.data
#train.data <- smote(diabetes ~ ., train.data, perc.over = 1)
test.data <- diabetes[-training.samples, ]
scaled_test.data <- scale(test.data[, 1:8])
test.data[, 1:8] <- scaled_test.data
#test.data <- na.omit(test.data)

# Handle NA's in training set
#mice <- complete(mice(subset(train.data, select = -c(triceps, insulin)), method='rf', seed = 123))
#mice <- complete(mice(train.data, method='rf', seed = 123))

#train.data$glucose <- mice$glucose

```

```

#train.data$pressure <- mice$pressure
#train.data$mass <- mice$mass

#train.data <- na.omit(train.data)

#mice.triceps <- complete(mice(subset(train.data, select = -insulin), method='rf', seed = 123))
#train.data$triceps <- mice.triceps$triceps

#mice.insulin <- complete(mice(train.data, method='rf', seed = 123))
#train.data$insulin <- mice.insulin$insulin

# Train
X.train <- model.matrix(diabetes~., data = train.data)[,-1]
#X.train <- scale(X.train)
y.train <- train.data$diabetes

# Test
X.test <- model.matrix(diabetes ~., data = test.data)[,-1]
#X.test <- scale(X.test)
y.test <- test.data$diabetes

```

Logistic Regression

```

set.seed(123)
log.model <- glm(diabetes ~., data = train.data, family = "binomial")

# Make predictions
probabilities <- predict(log.model, newdata = test.data, type = "response")
predicted.classes.log <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
# mean(predicted.classes == y.test)
log.conf <- confusionMatrix(predicted.classes.log, y.test, positive = "1")
log.conf

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 61 10
##           1  6 21
##
##               Accuracy : 0.8367
##               95% CI : (0.7484, 0.9037)
##       No Information Rate : 0.6837
##       P-Value [Acc > NIR] : 0.0004538
##
##               Kappa : 0.609
##
##  Mcnemar's Test P-Value : 0.4532547
##
##               Sensitivity : 0.6774
##               Specificity : 0.9104

```

```
##          Pos Pred Value : 0.7778
##          Neg Pred Value : 0.8592
##          Prevalence : 0.3163
##          Detection Rate : 0.2143
##          Detection Prevalence : 0.2755
##          Balanced Accuracy : 0.7939
##
##          'Positive' Class : 1
##
```

LASSO

```
set.seed(123)
cv.lasso.model <- cv.glmnet(X.train, y.train, alpha = 1, family = "binomial",
                           intercept = T)
#plot(cv.lasso.model)
cbind(coef(cv.lasso.model, s = cv.lasso.model$lambda.min), coef(cv.lasso.model, s = cv.lasso.model$lambda.1se))

## 9 x 2 sparse Matrix of class "dgCMatrix"
##          s1          s1
## (Intercept) -0.9087548 -0.79367862
## pregnant    0.2500270  0.04713706
## glucose     1.0571306  0.80577102
## pressure    .          .
## triceps     0.0626410  .
## insulin     .          .
## mass        0.4184043  0.19426613
## pedigree    0.3622318  0.10361089
## age         0.2235214  0.17759602

# Make predictions
probabilities <- predict(cv.lasso.model, newx = X.test, s = cv.lasso.model$lambda.min, type = "response")
predicted.classes.lasso.min <- as.factor(ifelse(probabilities > 0.5, 1, 0))

probabilities <- predict(cv.lasso.model, newx = X.test, s = cv.lasso.model$lambda.1se, type = "response")
predicted.classes.lasso.1se <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
#mean(predicted.classes.lasso.min == y.test)
lasso.min.conf <- confusionMatrix(predicted.classes.lasso.min, y.test, positive = "1")
lasso.min.conf

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 61 10
##          1  6 21
##
##          Accuracy : 0.8367
##          95% CI : (0.7484, 0.9037)
##          No Information Rate : 0.6837
##          P-Value [Acc > NIR] : 0.0004538
##
##          Kappa : 0.609
```

```
##
## McNemar's Test P-Value : 0.4532547
##
##          Sensitivity : 0.6774
##          Specificity : 0.9104
##          Pos Pred Value : 0.7778
##          Neg Pred Value : 0.8592
##          Prevalence : 0.3163
##          Detection Rate : 0.2143
##          Detection Prevalence : 0.2755
##          Balanced Accuracy : 0.7939
##
##          'Positive' Class : 1
##
lasso.1se.conf <- confusionMatrix(predicted.classes.lasso.1se, y.test, positive = "1")
lasso.1se.conf
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0  1
##          0 62 12
##          1  5 19
##
##          Accuracy : 0.8265
##          95% CI : (0.7369, 0.8956)
##          No Information Rate : 0.6837
##          P-Value [Acc > NIR] : 0.001064
##
##          Kappa : 0.573
##
## McNemar's Test P-Value : 0.145610
##
##          Sensitivity : 0.6129
##          Specificity : 0.9254
##          Pos Pred Value : 0.7917
##          Neg Pred Value : 0.8378
##          Prevalence : 0.3163
##          Detection Rate : 0.1939
##          Detection Prevalence : 0.2449
##          Balanced Accuracy : 0.7691
##
##          'Positive' Class : 1
##
```

in data: set.seed(123) and train 0.75 all others see(123) i like it a lot. min and 1se very sparse and good pred, just not much difference in amount of sparsity and amount of FN

in data: set.seed(2) and train 0.75 all others see(123) i like it a lot. 1se sparse and more FN but same pred

in data: set.seed(222) and train 0.75 all others see(123) pretty good similar too above

in data: set.seed(6) and train 0.75 all others see(123) gives very sparse 1se lasso and okay pred

in data: set.seed(13) and train 0.75 all others see(123) pretty nice! sparse 1se, but pred is same

in data: set.seed(42) and train 0.75 all others see(123) pretty nice! very very sparse 1se, more FN in 1se

some things we see:

- the more conservative model (1se) usually has more FN —> it makes the safe/conservative choice of going for the class that has 2/3 of observations

Ridge

```
set.seed(123)
cv.ridge.model <- cv.glmnet(X.train, y.train, alpha = 0, family = "binomial", intercept = T)
#plot(cv.ridge)

# Make predictions
probabilities <- predict(cv.ridge.model, newx = X.test, s = cv.ridge.model$lambda.min, type = "response")
predicted.classes.ridge <- as.factor(ifelse(probabilities > 0.5, 1, 0))

# Accuracy
#mean(predicted.classes.ridge == y.test)
ridge.conf <- confusionMatrix(as.factor(predicted.classes.ridge), y.test, positive = "1")
ridge.conf

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##          0 62 10
##          1  5 21
##
##              Accuracy : 0.8469
##              95% CI : (0.7601, 0.9117)
##      No Information Rate : 0.6837
##      P-Value [Acc > NIR] : 0.0001806
##
##              Kappa : 0.6301
##
##  Mcnemar's Test P-Value : 0.3016996
##
##              Sensitivity : 0.6774
##              Specificity : 0.9254
##      Pos Pred Value : 0.8077
##      Neg Pred Value : 0.8611
##              Prevalence : 0.3163
##      Detection Rate : 0.2143
##      Detection Prevalence : 0.2653
##      Balanced Accuracy : 0.8014
##
##      'Positive' Class : 1
##
```

Elastic Net

```
set.seed(123)

# CV and tuning grid
```

```

myFolds <- createFolds(y.train, k = 10, list = TRUE)
myControl <- trainControl(index = myFolds)

#cvControl <- trainControl(method = "cv", number = 10)
tuneGrid <- expand.grid(alpha = seq(0, 1, by = 0.05), lambda = 10^seq(1, -3, length=100))

# Train model
elasticnet.model <- train(X.train, y.train, method = "glmnet", trControl = myControl,
                          tuneGrid = tuneGrid, intercept = T, family = "binomial")

# Make predictions
probabilities <- predict(elasticnet.model, newdata = X.test, type = "prob") # "raw" already outputs pre
predicted.classes.elasticnet <- ifelse(probabilities$`1` > 0.5, 1, 0)

# Accuracy
# mean(predicted.classes == y.test)
elasticnet.conf <- confusionMatrix(as.factor(predicted.classes.elasticnet), y.test, positive = "1")
elasticnet.conf

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 63 11
##           1  4 20
##
##           Accuracy : 0.8469
##           95% CI : (0.7601, 0.9117)
##    No Information Rate : 0.6837
##    P-Value [Acc > NIR] : 0.0001806
##
##           Kappa : 0.6233
##
##  McNemar's Test P-Value : 0.1213353
##
##           Sensitivity : 0.6452
##           Specificity : 0.9403
##    Pos Pred Value : 0.8333
##    Neg Pred Value : 0.8514
##           Prevalence : 0.3163
##    Detection Rate : 0.2041
##    Detection Prevalence : 0.2449
##    Balanced Accuracy : 0.7927
##
##           'Positive' Class : 1
##
# Output coefficients
coef(elasticnet.model$finalModel, s = elasticnet.model$bestTune$lambda)


## 9 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -0.783644314
## pregnant    0.124519132
## glucose      0.616288277

```

```
## pressure      0.003625147
## triceps       0.073458775
## insulin       0.067361284
## mass          0.200698201
## pedigree      0.158043114
## age           0.188599082

# Filter data to only include specific alpha values
filtered_data <- subset(elasticnet.model$results, alpha %in% seq(0, 1, by=0.1))

# Create the plot
ggplot(filtered_data, aes(x=lambda, y=Accuracy, color=factor(alpha))) +
  geom_point() +
  geom_line() +
  scale_x_log10() +
  scale_color_discrete(name = "Alpha") +
  labs(x = "Lambda", y = "Accuracy") +
  theme_minimal()
```




```
library(ggcorrplot)
# Ensure that all remaining columns after subsetting are numeric
numeric_data <- subset(diabetes, select = -c(diabetes))
corr <- round(cor(numeric_data), 1)
ggcorrplot(corr,
  type = "lower",
  lab = TRUE,
  lab_size = 3,
```



```

colors = c("red", "white", "cyan4"),
title = "Correlogram of Diabetes Dataset",
ggtheme = theme_bw())

```




```

result <- check_collinearity(log.model)
plot(result) + ggtitle("")

```


```
## Variable 'Component' is not in your data frame :/
```

High collinearity (VIF) may inflate parameter uncertainty



Outliers

```
result <- check_outliers(diabetes, method = "mahalanobis")  
plot(result, type = "dots")
```



```
result
```

```
## 13 outliers detected: cases 2, 4, 5, 58, 78, 87, 111, 120, 185, 226,
##   236, 251, 302.
## - Based on the following method and threshold: mahalanobis (30).
## - For variables: pregnant, glucose, pressure, triceps, insulin, mass,
##   pedigree, age.
```

```
as.vector(result)
```

```
## [1] FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [37] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [49] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
## [61] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [73] FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE
## [85] FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [97] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [109] FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [121] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [133] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [145] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [169] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [181] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
## [193] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
## [205] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [217] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE
## [229] FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
## [241] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
## [253] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [265] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [277] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [289] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [301] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [313] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [325] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [337] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [349] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [361] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [373] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [385] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

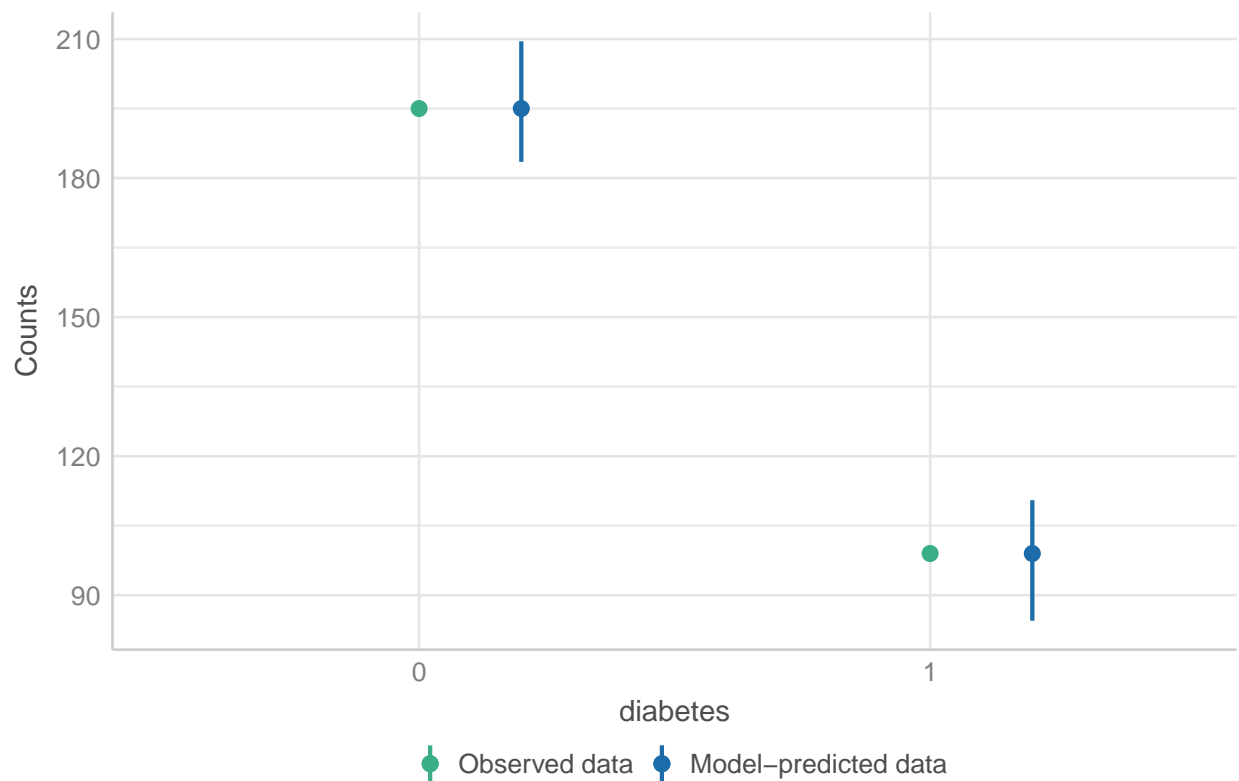
Check model

```
check <- check_model(log.model, panel = F)
plot(check)
```

```
## $PP_CHECK
```


Posterior Predictive Check

Model-predicted intervals should include observed data points



```
##
## $BINNED_RESID
```


Binned Residuals
Points should be within error bounds



\$OUTLIERS

Influential Observations


Points should be inside the contour lines



\$VIF

Collinearity

High collinearity (VIF) may inflate parameter uncertainty




##

\$QQ

Normality of Residuals

Dots should fall along the line



```
binmed_residuals(log.model)
```

```
## Warning: About 88% of the residuals are inside the error bounds (~95% or higher would be good).
```

Calculate model metrics

```
# Calculate metrics for each model
metrics.log <- calc_metrics(predicted.classes.log, y.test)
metrics.lasso.min <- calc_metrics(predicted.classes.lasso.min, y.test)
metrics.lasso.1se <- calc_metrics(predicted.classes.lasso.1se, y.test)
metrics.ridge <- calc_metrics(predicted.classes.ridge, y.test)
metrics.elasticnet <- calc_metrics(predicted.classes.elasticnet, y.test)

# Store in a dataframe
df_metrics <- as.data.frame(matrix(c(metrics.log, metrics.lasso.min, metrics.lasso.1se, metrics.ridge, metrics.elasticnet),
  nrow = 5, ncol = 5))
colnames(df_metrics) <- c("Logistic", "Lasso_min", "Lasso_1se", "Ridge", "ElasticNet")
rownames(df_metrics) <- c("AUC", "Recall", "Precision", "F1", "Accuracy")
df_metrics
```

```
##           Logistic Lasso_min Lasso_1se      Ridge ElasticNet
## AUC          0.7939336 0.7939336 0.7691382 0.8013962 0.7927299
## Recall       0.6774194 0.6774194 0.6129032 0.6774194 0.6451613
## Precision    0.7777778 0.7777778 0.7916667 0.8076923 0.8333333
## F1           0.7241379 0.7241379 0.6909091 0.7368421 0.7272727
## Accuracy     0.8367347 0.8367347 0.8265306 0.8469388 0.8469388
```


Latex Table Metrics

```
metrics_latex <- df_metrics[-1, ] %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), digits = 3, escape = FALSE,
        col.names = c("Logistic", "Lasso  $\lambda_{\min}$ ",
                      "Lasso  $\lambda_{1se}$ ", "Ridge", "ElasticNet")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Models" = 5), bold = TRUE) %>%
  row_spec(nrow(df_metrics) - 2, extra_latex_after = "\\midrule[.08em]") %>%
  row_spec(nrow(df_metrics) - 1, bold = TRUE)

print(metrics_latex)
```

	Models				
	Logistic	Lasso λ_{\min}	Lasso λ_{1se}	Ridge	ElasticNet
Recall	0.677	0.677	0.613	0.677	0.645
Precision	0.778	0.778	0.792	0.808	0.833
F1	0.724	0.724	0.691	0.737	0.727
Accuracy	0.837	0.837	0.827	0.847	0.847

Data Frame of all model coefficients

```
coef.log <- coef(log.model)
coef.lasso.min <- (coef(cv.lasso.model, s = cv.lasso.model$lambda.min))
coef.lasso.1se <- coef(cv.lasso.model, s = cv.lasso.model$lambda.1se)
coef.ridge <- coef(cv.ridge.model)
coef.elasticnet <- coef(elasticnet.model$finalModel, s = elasticnet.model$bestTune$lambda)

# Create a data frame to store coefficients
df_coef <- data.frame(
  Logistic = as.vector(coef.log),
  Lasso_min = as.vector(as.matrix(coef.lasso.min)),
  Lasso_1se = as.vector(as.matrix(coef.lasso.1se)),
  Ridge = as.vector(as.matrix(coef.ridge)),
  ElasticNet = as.vector(as.matrix(coef.elasticnet))
)
rownames(df_coef) <- names(coef.log)

options(scipen = 999)
df_coef_sparse <- round(df_coef, 3)
df_coef_sparse[] <- apply(df_coef_sparse, 2, function(x) ifelse(x == 0, '.', x))

df_coef_sparse

##           Logistic Lasso_min Lasso_1se Ridge ElasticNet
## (Intercept) -0.976   -0.909   -0.794 -0.777   -0.784
## pregnant    0.348     0.25    0.047  0.157     0.125
```

```
## glucose      1.197      1.057      0.806  0.437      0.616
## pressure    -0.008        .        .   0.084      0.004
## triceps      0.099      0.063        .   0.121      0.073
## insulin     -0.011        .        .   0.157      0.067
## mass         0.527      0.418      0.194  0.181      0.201
## pedigree     0.489      0.362      0.104  0.177      0.158
## age          0.24       0.224      0.178  0.186      0.189
```

Latex Table Coef

```
coef_latex <- df_coef_sparse %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Logistic", "Lasso  $\lambda_{\min}$ ",
                      "Lasso  $\lambda_{1se}$ ", "Ridge", "ElasticNet")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Models" = 5), bold = TRUE)

print(coef_latex)
```

	Models				
	Logistic	Lasso λ_{\min}	Lasso λ_{1se}	Ridge	ElasticNet
(Intercept)	-0.976	-0.909	-0.794	-0.777	-0.784
pregnant	0.348	0.25	0.047	0.157	0.125
glucose	1.197	1.057	0.806	0.437	0.616
pressure	-0.008	.	.	0.084	0.004
triceps	0.099	0.063	.	0.121	0.073
insulin	-0.011	.	.	0.157	0.067
mass	0.527	0.418	0.194	0.181	0.201
pedigree	0.489	0.362	0.104	0.177	0.158
age	0.24	0.224	0.178	0.186	0.189

Confusion matrix latex - log.model

```
# Convert the table to a data frame
log.conf.df <- as.data.frame(log.conf$table)


# Replace the numeric labels with text labels
log.conf.df$Prediction <- factor(log.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
log.conf.df$Reference <- factor(log.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))

# Rename columns as required by cums
names(log.conf.df) <- c('Target', 'Prediction', 'N')
# Plot the confusion matrix using cums
plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
  add_row_percentages = FALSE, add_col_percentages = FALSE, rotate_y_text = TRUE,
```

```
place_x_axis_above = TRUE, class_order = c("Positive", "Negative"))
```

```
## Warning in plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
## add_row_percentages = FALSE, : 'ggimage' is missing. Will not plot arrows and
## zero-shading.
```

```
## Warning in plot_confusion_matrix(log.conf.df, add_normalized = FALSE,
## add_row_percentages = FALSE, : 'rsvg' is missing. Will not plot arrows and
## zero-shading.
```



```
### Confusion matrix latex - lasso.min.model
```

```
# For lasso.min.conf
```

```
lasso.min.conf.df <- as.data.frame(lasso.min.conf$table)
```

```
lasso.min.conf.df$Prediction <- factor(lasso.min.conf.df$Prediction, levels = c("0", "1"), labels = c("N", "P"))
```


```
lasso.min.conf.df$Reference <- factor(lasso.min.conf.df$Reference, levels = c("0", "1"), labels = c("N", "P"))
```

```
names(lasso.min.conf.df) <- c('Target', 'Prediction', 'N')
```

```
plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_p
```

```
## Warning in plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.
```

```
## Warning in plot_confusion_matrix(lasso.min.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - lasso.1se.model

```
lasso.1se.conf.df <- as.data.frame(lasso.1se.conf$table)
lasso.1se.conf.df$Prediction <- factor(lasso.1se.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
lasso.1se.conf.df$Reference <- factor(lasso.1se.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))
names(lasso.1se.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)

## Warning in plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.

## Warning in plot_confusion_matrix(lasso.1se.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - ridge.model

```
# For ridge.conf
ridge.conf.df <- as.data.frame(ridge.conf$table)
ridge.conf.df$Prediction <- factor(ridge.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
ridge.conf.df$Reference <- factor(ridge.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))
names(ridge.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)

## Warning in plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.

## Warning in plot_confusion_matrix(ridge.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```




Confusion matrix latex - elasticnet.model

```
# For elasticnet.conf
elasticnet.conf.df <- as.data.frame(elasticnet.conf$table)
elasticnet.conf.df$Prediction <- factor(elasticnet.conf.df$Prediction, levels = c("0", "1"), labels = c("Negative", "Positive"))
elasticnet.conf.df$Reference <- factor(elasticnet.conf.df$Reference, levels = c("0", "1"), labels = c("Negative", "Positive"))
names(elasticnet.conf.df) <- c('Target', 'Prediction', 'N')
plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, add_row_percentages = FALSE, add_col_percentages = FALSE)


## Warning in plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, :
## 'ggimage' is missing. Will not plot arrows and zero-shading.

## Warning in plot_confusion_matrix(elasticnet.conf.df, add_normalized = FALSE, :
## 'rsvg' is missing. Will not plot arrows and zero-shading.
```

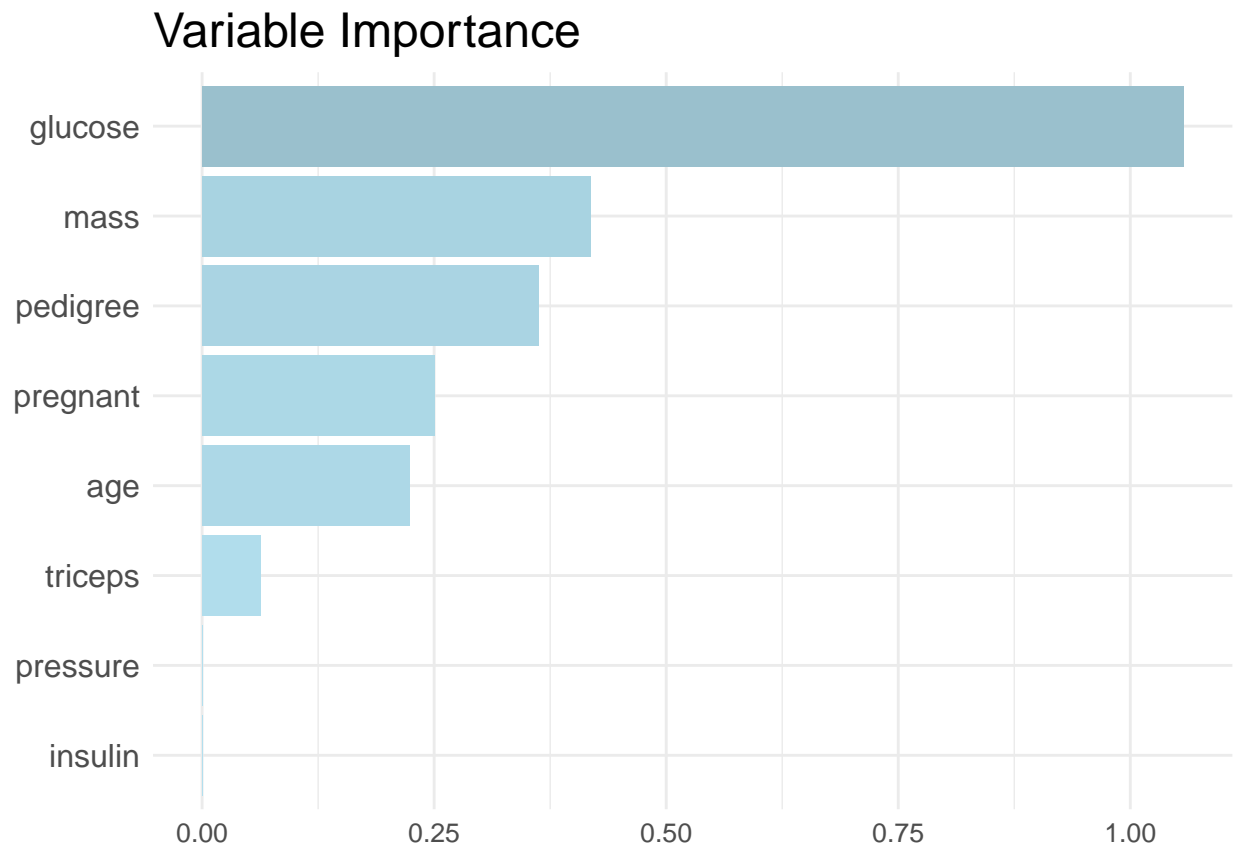


Variable Importance not working great because font size

```
#png(filename = "varimp_log.png", width = 3000, height = 3000)  
importance_plot("Logistic")
```




```
#dev.off()
#png(filename = "varimp_lasso_min.png", width = 3000, height = 3000)
importance_plot("Lasso_min")
```


```
#dev.off()  
#png(filename = "varimp_lasso_1se.png", width = 3000, height = 3000)  
importance_plot("Lasso_1se")
```

Variable Importance

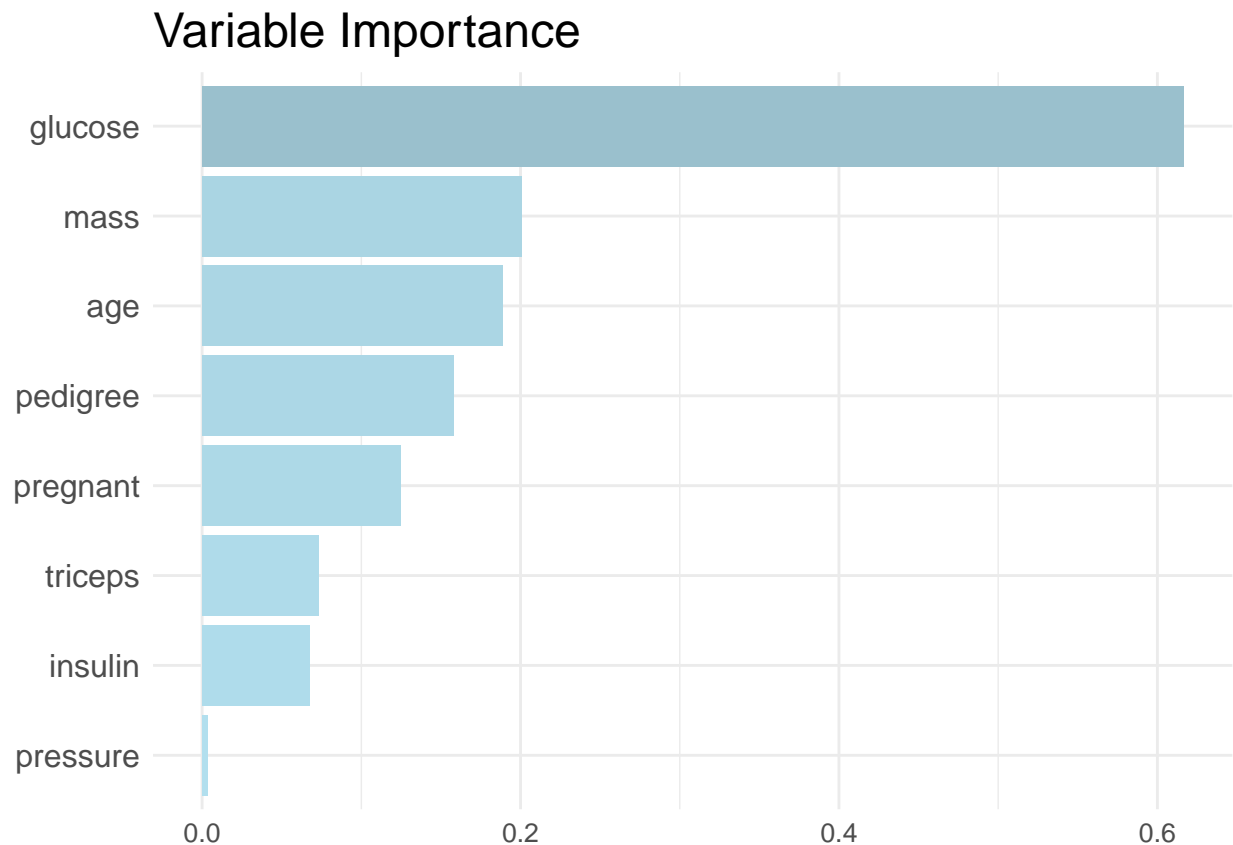


```
#dev.off()  
#png(filename = "varimp_ridge.png", width = 3000, height = 3000)  
importance_plot("Ridge")
```

Variable Importance



```
#dev.off()  
#png(filename = "varimp_elasticnet.png", width = 3000, height = 3000)  
importance_plot("ElasticNet")
```



```
#dev.off()
```

Perc missing

```
library(ggplot2)
```

```
# Count of missing values per column
```

```
count_na <- c(pregnant = 0, glucose = 5, pressure = 35, triceps = 227, insulin = 374, mass = 11, pedigree = 11)
```

```
# Total number of rows in the data
```

```
n_rows <- nrow(diabetes)
```

```
# Convert to data frame and calculate percentages
```

```
df_na <- data.frame(Column = names(count_na), Count = as.numeric(count_na))
```

```
df_na$Percentage <- (df_na$Count / n_rows) * 100
```

```
# Plot
```

```
ggplot(df_na, aes(x = reorder(Column, -Percentage), y = Percentage, fill = Percentage)) +
```

```
  geom_bar(stat = "identity") +
```

```
  scale_fill_gradient(low = "lightblue1", high = "lightblue3") +
```

```
  theme_minimal() +
```

```
  theme(
```

```
    axis.text.y = element_text(size = 12),
```

```
    axis.text.x = element_text(size = 12),
```

```
    title = element_text(size = 15),
```

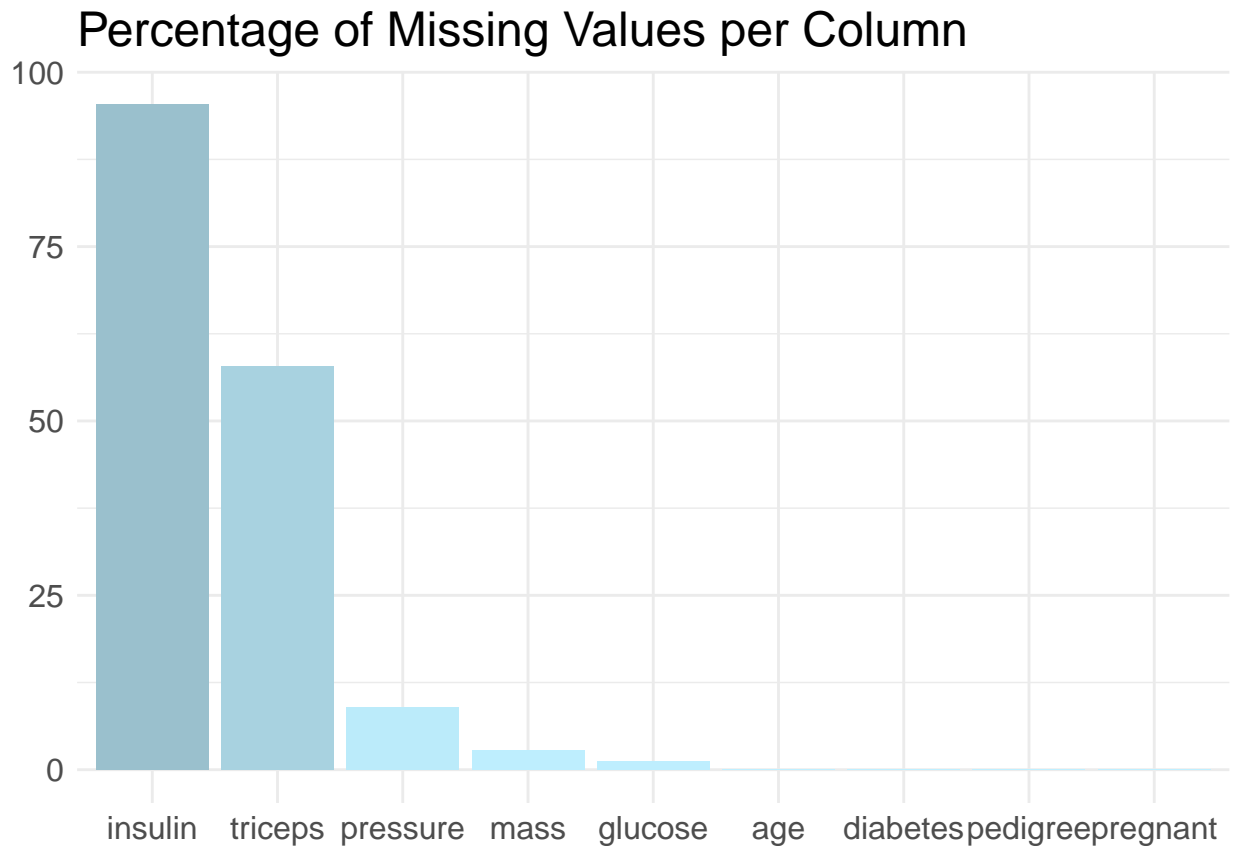
```
    axis.title.y = element_blank(),
```

```
    axis.title.x = element_blank(),
```

```

legend.position = "none"
) +
ggtitle("Percentage of Missing Values per Column")

```



log model summary table latex

```

# Extract coefficients from the model summary
coef_summary <- summary(log.model)$coefficients

# Convert the matrix to a data frame for kable()
coef_summary <- round(as.data.frame(coef_summary)[, c(1, 4)], 3)
colnames(coef_summary) <- c("Estimate", "p-value")

# coef_summary$p-value <- cell_spec(coef_summary$p-value, "latex",
#                                   bold = ifelse(coef_summary$p-value < 0.05, TRUE, FALSE))

# # Create the LaTeX table
# kable(coef_summary, "latex", booktabs = TRUE, align = "c", escape = F) %>%
# kable_styling(latex_options = c("striped", "scale_down", "hold_position")) %>%
# row_spec(0, bold = TRUE)

# Add stars based on significance level and make significant p-values bold
coef_summary$p-value <- sapply(coef_summary$p-value, function(p) {
  stars <- ifelse(p < 0.001, "***", ifelse(p < 0.01, "**", ifelse(p < 0.05, "*", "")))
  value_with_stars <- paste0(round(p, 3), stars)
  cell_spec(value_with_stars, "latex", bold = p < 0.05)
})

```

```
# Create the LaTeX table
```

```
kable(coef_summary, "latex", booktabs = TRUE, align = "c", escape = F) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position")) %>%
  row_spec(0, bold = TRUE)
```

	Estimate	p-value
(Intercept)	-0.976	0***
pregnant	0.348	0.087
glucose	1.197	0***
pressure	-0.008	0.963
triceps	0.099	0.622
insulin	-0.011	0.95
mass	0.527	0.016*
pedigree	0.489	0.004**
age	0.240	0.251

```
coef_latex <- coef_summary %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Estimate", "p value")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  #add_header_above(c(" " = 1, "Logistic Regression" = 2), bold = TRUE)

print(coef_latex)
```

coef latex per model - lasso both

```
coef_latex <- df_coef_sparse[, 2:3] %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Lasso  $\\lambda_{\\text{min}}$ ", "Lasso  $\\lambda_{\\text{1se}}$ ")) %>%
```

	Estimate	p value
(Intercept)	-0.976	0***
pregnant	0.348	0.087
glucose	1.197	0***
pressure	-0.008	0.963
triceps	0.099	0.622
insulin	-0.011	0.95
mass	0.527	0.016*
pedigree	0.489	0.004**
age	0.240	0.251

```
kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Estimates" = 2), bold = TRUE)

print(coef_latex)
```

coef latex per model - log, ridge

```
coef_latex <- df_coef_sparse[, c(1,4)] %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Logistic", "Ridge")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Estimates" = 2), bold = TRUE)

print(coef_latex)
```

	Estimates	
	Lasso λ_{\min}	Lasso λ_{1se}
(Intercept)	-0.909	-0.794
pregnant	0.25	0.047
glucose	1.057	0.806
pressure	.	.
triceps	0.063	.
insulin	.	.
mass	0.418	0.194
pedigree	0.362	0.104
age	0.224	0.178

coef latex per model - log, elastic net

```
coef_latex <- df_coef_sparse[, c(1,5)] %>%
  kable(format = "latex", booktabs = TRUE, align = c('c'), escape = FALSE,
        col.names = c("Logistic", "ElasticNet")) %>%
  kable_styling(latex_options = c("striped", "scale_down", "hold_position"), position = "center") %>%
  row_spec(0, bold = TRUE) %>%
  column_spec(1, bold = TRUE) %>%
  add_header_above(c(" " = 1, "Estimates" = 2), bold = TRUE)

print(coef_latex)

# Plot the coefficient paths with thicker lines
plot(cv.lasso.model$glmnet.fit, xvar = "lambda", label = TRUE, cex.axis = 1.5, lwd = 2)


# Add vertical lines for lambda.min and lambda.1se with thicker lines
abline(v = log(cv.lasso.model$lambda.min), col = "red", lty = 2, lwd = 2)
abline(v = log(cv.lasso.model$lambda.1se), col = "blue", lty = 2, lwd = 2)

# Add a legend with LaTeX-style text and increased font size, and thicker lines
legend("topright",
      legend = expression(lambda["min"], lambda["1se"]),
```


	Estimates	
	Logistic	Ridge
(Intercept)	-0.976	-0.777
pregnant	0.348	0.157
glucose	1.197	0.437
pressure	-0.008	0.084
triceps	0.099	0.121
insulin	-0.011	0.157
mass	0.527	0.181
pedigree	0.489	0.177
age	0.24	0.186

```
col = c("red", "blue"),
lty = 2,
cex = 1.5,
lwd = 0.5, bty = "n")
```

	Estimates	
	Logistic	ElasticNet
(Intercept)	-0.976	-0.784
pregnant	0.348	0.125
glucose	1.197	0.616
pressure	-0.008	0.004
triceps	0.099	0.073
insulin	-0.011	0.067
mass	0.527	0.201
pedigree	0.489	0.158
age	0.24	0.189



```

p value Acc > NIR
# Number of trials
n <- 98

# Number of successes (correctly classified instances)
# Accuracy is 0.8367, so the number of successes would be 0.8367 * 392
n_success <- round(0.8367 * n)

# No Information Rate
nir <- 0.6837

# Perform binomial test
binom_test_result <- binom.test(n_success, n, p = nir, alternative = "greater")

# Print the result
print(binom_test_result)

##
## Exact binomial test
##
## data: n_success and n
## number of successes = 82, number of trials = 98, p-value = 0.0004546
## alternative hypothesis: true probability of success is greater than 0.6837
## 95 percent confidence interval:
##  0.7626479 1.0000000
## sample estimates:
## probability of success
##           0.8367347
log.conf$overall["Accuracy"]

Accuracy 0.8367347
log.conf$overall["AccuracyPValue"]

AccuracyPValue 0.0004537699
tibble(Accuracy = log.conf$overall["Accuracy"],
       `p-value [Acc > NIR]` = log.conf$overall["AccuracyPValue"])

```

A tibble: 1 x 2

```

Accuracy p-value [Acc > NIR] 1 0.837 0.000454
result_table <- tibble(Accuracy = log.conf$overall["Accuracy"],
                      `p-value [Acc > NIR]` = log.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table$`p-value [Acc > NIR]` <- round(result_table$`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

Accuracy	p-value [Acc > NIR]
----------	---------------------

0.8367

0.0005

```
result_table <- tibble(Accuracy = log.conf$overall["Accuracy"],
  `p-value [Acc > NIR]` = log.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table$`p-value [Acc > NIR]` <- round(result_table$`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Accuracy	p-value [Acc > NIR]
----------	---------------------

0.8367

0.0005

```
result_table <- tibble(Accuracy = ridge.conf$overall["Accuracy"],
  `p-value [Acc > NIR]` = ridge.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table$`p-value [Acc > NIR]` <- round(result_table$`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Accuracy	p-value [Acc > NIR]
----------	---------------------

0.8469

0.0002

```
#c("Lasso  $\lambda_{\min}$ ", "Lasso  $\lambda_{1se}$ ")
result_table <- tibble(Accuracy = lasso.min.conf$overall["Accuracy"],
  `p-value [Acc > NIR]` = lasso.min.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table`p-value [Acc > NIR]` <- round(result_table`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Accuracy	p-value [Acc > NIR]
0.8367	0.0005

```
result_table <- tibble(Accuracy = lasso.1se.conf$overall["Accuracy"],
  `p-value [Acc > NIR]` = lasso.1se.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table`p-value [Acc > NIR]` <- round(result_table`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Accuracy	p-value [Acc > NIR]
0.8265	0.0011

```
result_table <- tibble(Accuracy = elasticnet.conf$overall["Accuracy"],
  `p-value [Acc > NIR]` = elasticnet.conf$overall["AccuracyPValue"])

# Round values for better formatting, if needed
result_table$Accuracy <- round(result_table$Accuracy, 4)
result_table`p-value [Acc > NIR]` <- round(result_table`p-value [Acc > NIR]`, 4)

# Generate LaTeX table
kable(result_table, "latex", booktabs = TRUE, align = "c") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Accuracy	p-value [Acc > NIR]
0.8469	0.0002